

Giannis Karamanolakis

Department of Computer Science, Columbia University
New York, NY 10027
Office: Mudd 416, Data Science Institute (map)

WWW: <https://gkaramanolakis.github.io>
E-mail: gkaraman@cs.columbia.edu
Phone: +1 (917) 331-8844

Research Overview and Interests

My research has been focusing on **minimally-supervised learning** frameworks to address real-world Natural Language Processing applications without the costly requirement of large-scale labeled data. My work leverages alternative types of expert knowledge as supervision for deep neural networks across domains, tasks, and languages, and has been shipped into real-world operational systems with an immediate impact on humans.

Research interests: Machine Learning, Information Extraction, Natural Language Processing, Deep Learning, Interactive Machine Teaching, Semi/Weakly-Supervised Learning, Cross-lingual Learning, Question Answering, Semantic Parsing.

Professional and Research Experience

5. **Columbia University**, Department of Computer Science, New York, NY, United States. 09/2017 - now
Graduate Research Assistant
Advisors: Luis Gravano, Daniel Hsu..
 - Developing minimally-supervised learning algorithms for text mining (WNUT'19, LLD'19, EMNLP'19/20, LOUHI'20).
 - Deploying systems for "Information Extraction from Social Media for Actionable Inferences in Public Health" (NSF).
 - Actively participating in the Machine Learning group, NLP group, and Database Research group.
4. **Microsoft Research AI**, Redmond, WA, United States. 05/2020 - 08/2020
Research Scientist Intern
Supervisors: Ahmed Awadallah, Subho Mukherjee, Guoqing Zheng (Language & Information Technologies team).
 - Developed ASTRA, a semi-supervised learning method for self-training with expert-defined rules (NAACL '21 paper).
 - Designed a benchmark on semi/weakly-supervised learning for natural language tasks (under submission).
3. **Amazon.com**, Seattle, WA, United States. 05/2019 - 12/2019
Applied Scientist / Machine Learning Scientist Intern
Supervisor: Xin Luna Dong (Product Graph team).
 - Developed TXtract, a neural network for large-scale knowledge extraction from product descriptions (ACL '20 paper).
 - Improved product understanding through 15% higher coverage in knowledge graph construction (KDD '20 paper).
2. **Behavioral Signals**, Los Angeles, United States (worked remotely from Athens, Greece). 04/2016 - 08/2017
Machine Learning Engineer
Supervisors: Shri Narayanan, Alexandros Potamianos.
 - Trained statistical models for recognizing core emotions from conversational speech.
 - Used speaker interaction patterns to improve the prediction of key performance indicators.
1. **NCSR "Demokritos"**, Computational Intelligence Laboratory, Athens, Greece. 11/2015 - 2/2016
Research Intern
Supervisors: Theodoros Giannakopoulos, Stavros Perantonis.
 - Developed audio-based classifiers for acoustic event detection and soundscape quality estimation.

Education

2. **Columbia University**, New York, NY, United States. 09/2017 - now (expected: 05/2022)
Ph.D. in Computer Science GPA: 4.26/4.00
 Advisors: Prof. Luis Gravano, Prof. Daniel Hsu.
1. **National Technical University of Athens**, Athens, Greece. 09/2011 - 07/2017
M.Eng. in Electrical & Computer Engineering (5-year study program, MSc equivalent) GPA: 9.26/10.00
 Advisor: Prof. Alexandros Potamianos.
 Dissertation: "Grounding Word Semantic Representations to Acoustic and Visual Perceptual Modalities".

Selected Publications

5. **Giannis Karamanolakis**, Subhabrata Mukherjee, Guoqing Zheng, and Ahmed Hassan Awadallah. **Self-Training with Weak Supervision**. In *NAACL*, 2021 (Oral).
4. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Cross-Lingual Text Classification with Minimal Resources By Transferring a Sparse Teacher**. In *Findings of EMNLP*, 2020.
3. **Giannis Karamanolakis**, Jun Ma, and Xin Luna Dong. **TXtract: Taxonomy-Aware Knowledge Extraction for Thousands of Product Categories**. In *ACL*, 2020 (Oral).
2. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Leveraging Just a Few Keywords for Fine-Grained Aspect Detection Through Weakly Supervised Co-Training**. In *EMNLP*, 2019 (Oral).
1. **Giannis Karamanolakis**, Elias Iosif, Athanasia Zlatintsi, Aggelos Pikrakis, and Alexandros Potamianos. **Audio-Based Distributional Representations of Meaning Using a Fusion of Feature Encodings**. In *INTERSPEECH*, 2016 (Oral).

Selected Fellowships, Honors, and Awards

7. **Leventis Foundation Scholarship**, New York, US 08/2019, 08/2020, 08/2021
6. **Gerondelis Foundation Scholarship**, New York, US 06/2018
5. **Joseph F. Traub Fellowship**, Columbia University, New York, US 04/2018
4. **Addressing the Opioid Epidemic, Design Challenge Award**, Columbia University, New York, US 12/2017
3. **Columbia Engineering Presidential Fellowship**, New York, US 08/2017 - 05/2022
2. **Eurobank EFG Scholarship**, Athens, Greece 06/2012
1. **Hellenic Mathematic Society "Euclides" & "Thales" Awards**, Sitia, Greece 02/2009, 03/2010

Technical Skills, Talks, and Academic Service

- **Programming:** Python (PyTorch, Tensorflow, Keras, Pyro, Huggingface, Spacy, Gensim), C++, Java, MATLAB, C.
- **Agile Development and Project Management:** Git, JIRA, Jenkins, Scrum, Kanban.
- **Invited Industry Talks:** Two Sigma (Dec. '20), RelationalAI (May '20), NYAS (Mar. '20), AT&T Labs (Dec. '19).
- **Conference/Workshop Talks:** NAACL '21, DeeLIO '20, ACL '20, EMNLP '19, RecSys '18, InterSpeech '16.
- **Reviewer Activity:** ICML '21-22 (expert reviewer), NeurIPS '19-21 (top 10% reviewer), ACL Rolling Review '21, ACL '19-21, EMNLP '21, Machine Learning Journal '21, TrueFact '20, TKDD journal '19, WWW '19, IEEE-TETCI '18.
- **Mentoring:** Designed research projects for 12 students (8 undergraduate, 4 graduate); 5 published papers.

All Publications

15. **Giannis Karamanolakis**, Subhabrata Mukherjee, Guoqing Zheng, and Ahmed Hassan Awadallah. **Self-Training with Weak Supervision**. In *NAACL*, 2021 (Oral).
14. Ivy Cao, Zizhou Liu, **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Quantifying the Effects of COVID-19 on Restaurant Reviews**. In *EMNLP 9th International Workshop on Natural Language Processing for Social Media*, 2021 (Oral).
13. Tejit Pabari, **Giannis Karamanolakis**, Max Mauerman, Beth Tellman, Mitchell Thomas, and Eugene Wu. **Flood Event Extraction from Newspaper Media in Bangladesh**. In *Book of Abstracts of the 8th International Conference on Water and Flood Management (ICWFM)*, 2021.
12. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Cross-Lingual Text Classification with Minimal Resources By Transferring a Sparse Teacher**. In *Findings of EMNLP*, 2020.
11. Ziyi Liu, **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Detecting Foodborne Illness Complaints in Multiple Languages Using English Annotations Only**. In *EMNLP Workshop on Health Text Mining and Information Analysis*, 2020 (Oral).
10. **Giannis Karamanolakis**, Jun Ma, and Xin Luna Dong. **TXtract: Taxonomy-Aware Knowledge Extraction for Thousands of Product Categories**. In *ACL*, 2020 (Oral).
9. Xin Luna Dong, Xiang He, Andrey Kan, Xian Li, Yan Liang, Jun Ma, Yifan Ethan Xu, Chenwei Zhang, Tong Zhao, Gabriel Blanco Saldana, Saurabh Deshpande, Alexandre Michetti Manduca, Jay Ren, Surender Pal Singh, Fan Xiao, Haw-Shiuan Chang, **Giannis Karamanolakis**, Yuning Mao, Yaqing Wang, Christos Faloutsos, Andrew McCallum, and Jiawei Han. **AutoKnow: Self-Driving Knowledge Collection for Products of Thousands of Types**. In *KDD*, 2020 (Oral).
8. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Leveraging Just a Few Keywords for Fine-Grained Aspect Detection Through Weakly Supervised Co-Training**. In *EMNLP*, 2019 (Oral).
7. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **weakly supervised attention networks for fine-grained opinion mining and public health**. In *EMNLP Workshop on Noisy User-generated Text*, 2019 (Oral).
6. **Giannis Karamanolakis**, Daniel Hsu, and Luis Gravano. **Training Neural Networks for Aspect Extraction Using Descriptive Keywords Only**. In *ICLR 2nd Workshop on Learning from Limited Labeled Data*, 2019.
5. Katy Ilonka Gero, **Giannis Karamanolakis**, and Lydia Chilton. **Transfer Learning for Style-Specific Text Generation**. In *NeurIPS Workshop on Machine Learning for Creativity and Design*, 2018.
4. **Giannis Karamanolakis**, Kevin Raji Cherian, Ananth Ravi Narayan, Jie Yuan, Da Tang, and Tony Jebara. **Item Recommendation with Variational Autoencoders and Heterogenous Priors**. In *RecSys Workshop on Deep Learning for Recommender Systems*, 2018 (Oral).
3. Georgios Paraskevopoulos, **Giannis Karamanolakis**, Elias Iosif, Aggelos Pikrakis, and Alexandros Potamianos. **Sensory-Aware Multimodal Fusion for Word Semantic Similarity Estimation**. In *EUSIPCO Multimodal Processing, Modeling, and Learning (Multi-Learn) Workshop*, 2017 (Oral).
2. **Giannis Karamanolakis**, Elias Iosif, Athanasia Zlatintsi, Aggelos Pikrakis, and Alexandros Potamianos. **Audio-based Distributional Semantic Models for Music Auto-tagging and Similarity Measurement**. In *EUSIPCO Multimodal Processing, Modeling, and Learning Workshop*, 2017 (Oral).
1. **Giannis Karamanolakis**, Elias Iosif, Athanasia Zlatintsi, Aggelos Pikrakis, and Alexandros Potamianos. **Audio-Based Distributional Representations of Meaning Using a Fusion of Feature Encodings**. In *INTERSPEECH*, 2016 (Oral).